

MUSTAFA SHOUKAT

Senior AI Engineer, Enterprise AI Architecture and Technical Leadership

mustafashoukat.ai@gmail.com linkedin.com/in/mustafashoukat github.com/Mustafa-Shoukat1 Riyadh, Saudi Arabia

Executive Synopsis

Prepared for Staff AI Engineer, Teradata, 23 September 2026

Profile

Mustafa Shoukat is an AI engineer from Pakistan, based in Riyadh, Saudi Arabia, who operates at Staff level scope: he owns the architecture of distributed, fault tolerant AI platforms end to end, not just individual features. He has architected and shipped more than 20 production systems for enterprise and government grade clients across Saudi Arabia, the UAE, the United States and the United Kingdom, taking AI from prototype to hardened, observable, day to day production inside regulated environments.

Career trajectory

- **Data Scientist, COMET Estimating.** ML, NLP and document processing pipelines hardened against messy real world construction data.
- **Generative AI Consultant, Aretec.** Production LLM and retrieval applications for business users, including fine tuning for accuracy and response speed.
- **Agentic AI Engineer, Watad (Navid), Riyadh.** Architecture owner and team lead for ByanRAG 2.2, the company's Arabic first enterprise RAG platform, and contributor to Yehia, Navid's Arabic LLM.
- **Independent AI Software Engineer.** Architects and ships production AI platforms for clients in four countries and advises organizations moving AI from pilot to operations.

Most relevant enterprise AI achievements

- **ByanRAG 2.2.** Multi tenant, Arabic first RAG platform used by two organizations and more than 150 users, with isolated vector collections per organization, five levels of role based access on every endpoint, full audit logging, and on premises deployment for government networks. Cited answers in 2 to 3 seconds on self hosted open source models.
- **Yehia R1.** Designed the Claude as judge reward pipeline (3C3H metric) and NLI evaluation guardrails that helped take Yehia R1 to first place on the AraGen Leaderboard (0.5B to 25B).
- **Global Holman.** Agentic government contracting platform: LangGraph and CrewAI workflows that read SAM.gov solicitations, analyze requirements, support bid decisions and draft PDF proposals.
- **VacantSeek 2.0.** B2B real estate SaaS on AWS ECS; led a production hardening program of more than 340 merged pull requests covering deployment, storage, security gates and monitoring.
- **Whoza.ai.** Live AI voice receptionist for UK tradespeople, built with a team of four on Twilio, FastAPI, React and Supabase.

Technical leadership

He leads and mentors the ByanRAG team of four (an AI engineer, an LLM engineer, a frontend engineer and a DevOps and QA engineer), sets the standards for testing, deployment and model integration, backed by a suite of more than 600 tests, and has trained more than 150 professionals in data science, ML and AI engineering. He contributes to LangChain, RAGAS and Unstructured, released the open source tool ytkb, and builds with MCP, including his own MCP server with Arabic tools.

Core stack

Python and Java, FastAPI, LangGraph, CrewAI and MCP, Qdrant, Milvus, FAISS and pgvector, PostgreSQL and MongoDB, Celery and Redis, Docker, AWS, GCP and Azure, Prometheus, Grafana and Loki.