

Mustafa Shoukat

Senior AI Engineer · Enterprise AI Architecture & Technical Leadership

mustafashoukat.ai@gmail.com · +92 309 3609261 · LinkedIn · GitHub

Prepared for: Staff AI Engineer — Teradata

This brief maps Mustafa Shoukat’s experience directly against the published requirements of Teradata’s Staff AI Engineer role. Each row states the requirement and the concrete, production evidence behind it.

REQUIREMENT-BY-REQUIREMENT FIT

Teradata requirement	Evidence from Mustafa’s work
Design & deliver highly scalable, next-gen AI products	Owned end-to-end architecture for ByanRAG 2.2 , a production multi-tenant RAG platform, and a LangGraph agentic government-contracting platform — both live, not prototypes, across enterprise/government clients in four countries.
Architect for scalability, reliability & performance	Distributed, fault-tolerant design with Celery/Redis task processing, org-isolated data, hybrid retrieval tuned for latency (RRF + CrossEncoder reranking), backed by a 600+ test suite and CI/CD .
RESTful APIs with security & data consistency	Built scalable platform APIs in FastAPI ; fail-closed HMAC webhook verification, atomic message deduplication, and reusable, provider-agnostic service layers.
SQL & NoSQL data modeling	Production use of PostgreSQL (async), MongoDB , Redis, pgvector, Supabase/Neon (SQLAlchemy/Alembic) — plus a Text-to-SQL RAG app translating plain English into queries over complex enterprise schemas.
AI/ML, agentic AI & LLMs	Core strength: LangGraph , CrewAI , MCP multi-agent orchestration with tool-calling and multi-step planning; hybrid RAG; a Claude-as-judge pipeline that reached #1 on the AraGen Leaderboard .
Distributed system patterns & debugging	Fault-tolerant Celery/Redis architecture; production observability via Prometheus/Grafana/Loki + LangSmith/Sentry used to resolve real bottlenecks (e.g. reranker latency).
Containerization / orchestration	Containerized services with Docker & Docker Compose , deployed behind nginx/TLS with CI/CD. (<i>Kubernetes: see note below.</i>)
Public cloud	Hands-on across AWS (S3, App Runner) , GCP , and Azure .
Programming languages (Java / Go / Python)	Python (primary) and Java ; also TypeScript, JavaScript, SQL, Rust.
AI productivity tools (e.g. GitHub Copilot)	Daily practitioner of AI-assisted development and modern LLM tooling.
Mentorship & influencing architecture	Set engineering standards across the team, mentored engineers, and trained 150+ professionals ; merged open-source PRs to LangChain , RAGAS , Unstructured .
Communication & independent delivery	Ships independently for clients in four countries; advises organizations moving AI from pilot to production.

THREE FLAGSHIP PROOFS

1. ByanRAG 2.2 — enterprise trust at scale

A multi-tenant, Arabic-first document-management and RAG platform with per-organization data isolation, a 5-level Casbin RBAC hierarchy enforced on every endpoint, a 14-panel admin console, full audit logging, and **on-premises / air-gapped deployment inside regulated government networks**. Exactly the enterprise-grade, security-first AI delivery a Staff role demands.

2. #1 on the AraGen Leaderboard — measurable model quality

A Claude-as-judge reward pipeline (3C3H metric) plus NLI-based evaluation guardrails — faithfulness/groundedness scoring, Recall@5, MRR, and prompt-injection defenses — that took **Yehia R1 to #1 (0.5B–25B)**. Evidence of rigor in evaluation and production readiness, not just model-building.

3. Agentic government-contracting platform — orchestration in production

LangGraph multi-agent workflows scraping SAM.gov by NAICS code, ingesting solicitation PDFs into a RAG pipeline, and automating requirement analysis, bid/no-bid decisions, and proposal generation over a provider-agnostic LLM layer — replacing manual analyst workflows.

NOTES FOR ACCURACY

- **Kubernetes:** production orchestration to date is **Docker / Docker Compose behind nginx/TLS**, not Kubernetes; the container and distributed-systems foundation transfers directly.
- **Messaging systems (Kafka, preferred):** event-driven experience is via **Celery/Redis, WebSockets, and SSE** rather than Kafka.
- **Experience depth: 5+ years of senior-scope AI engineering** with full architecture ownership across 20+ production systems — the strength is the level of ownership and production impact, which maps to Staff-level responsibilities.

SUMMARY

Mustafa meets or exceeds the core of this role — **AI platform architecture, RAG, agentic orchestration, distributed systems, production readiness, and technical leadership** — with live, enterprise-grade systems and a top-of-leaderboard model result to prove it. The few adjacencies (Kubernetes, Kafka) are quickly bridged by an engineer who already owns production distributed systems end-to-end.